

ADVANCED SPARK ASSIGNMENT 21.2

PROBLEM STATEMENT:

- 1) Join of two or more data sets is one of the most widely used operations you do with your data, but in distributed systems it can be a huge headache. In general, since your data are distributed among many nodes, they have to be shuffled before a join that causes significant network I/O and slow performance.
- 2) Fortunately, if you need to join a large table with relatively small tables you can avoid sending all data of the large table over the network. This type of join is called map-side join in Hadoop community. In other distributed systems, it is often called replicated or broadcast join. The fact table can be very large, while dimension tables are often quite small.
- 3) Let's use the following sample data (one fact and two-dimension tables):

// Fact table

```
val flights = sc.parallelize(List(
  ("SEA", "JFK", "DL", "418", "7:00"),
  ("SFO", "LAX", "AA", "1250", "7:05"),
  ("SFO", "JFK", "VX", "12", "7:05"),
  ("JFK", "LAX", "DL", "424", "7:10"),
  ("LAX", "SEA", "DL", "5737", "7:10")
))
```

// Dimension table/broadcast table

```
val airports = sc.parallelize(List(
  ("JFK", "John F. Kennedy International Airport", "New York", "NY"),
  ("LAX", "Los Angeles International Airport", "Los Angeles", "CA"),
  ("SEA", "Seattle-Tacoma International Airport", "Seattle", "WA"),
  ("SFO", "San Francisco International Airport", "San Francisco", "CA")))

```

// Dimension table/broadcast table

```
val airlines = sc.parallelize(List(
  ("AA", "American Airlines"),
  ("DL", "Delta Airlines"),
  ("VX", "Virgin America")))

```

We need to join the fact and dimension tables to get the following result:

San Francisco	Los Angeles	American Airlines	1250	7:05
New York	Los Angeles	Delta Airlines	424	7:10
Seattle	New York	Delta Airlines	418	7:00
San Francisco	New York	Virgin America	12	7:05
Los Angeles	Seattle	Delta Airlines	5737	7:10

//Importing the spark sql packages

```
import org.apache.spark.sql._
import sqlContext.implicits._
```

//putting the flights data into SPARK RDD

```
val flights = sc.parallelize(List(
  ("SEA", "JFK", "DL", "418", "7:00"),
```

```
("SFO", "LAX", "AA", "1250", "7:05"),
("SFO", "JFK", "VX", "12", "7:05"),
("JFK", "LAX", "DL", "424", "7:10"),
("LAX", "SEA", "DL", "5737", "7:10")
))
```

```
//putting the airports data into SPARK RDD
```

```
val airports = sc.parallelize(List(
  ("JFK", "John F. Kennedy International Airport", "New York", "NY"),
  ("LAX", "Los Angeles International Airport", "Los Angeles", "CA"),
  ("SEA", "Seattle-Tacoma International Airport", "Seattle", "WA"),
  ("SFO", "San Francisco International Airport", "San Francisco", "CA")
))
```

```
//putting the airlines data into SPARK RDD
```

```
val airlines = sc.parallelize(List(
  ("AA", "American Airlines"),
  ("DL", "Delta Airlines"),
  ("VX", "Virgin America")
))
```

```
-----
//Conversion of flights RDD into Dataset
```

```
val fds = flights.toDS()
```

```
//Conversion of Dataset to SPARK Dataframe
```

```
val firstdf = fds.toDF()
```

```
//Naming the Columns of SPARK Dataframe into putting it into RDD
```

```
val columnnamef =
Seq("sourceabbrev", "airportabbrev", "airlineabbrev", "track", "depttime")
```

```
//Renaming the columns of SPARK Dataframe with the help of SPARK RDD passed
as an argument
```

```
val firstddf = firstdf.toDF(columnnamef: _*)
```

```
-----
//Conversion of airports RDD into Dataset
```

```
val apds = airports.toDS()
```

```
//Conversion of airport Dataset to SPARK Dataframe
```

```
val apdf = apds.toDF()
```

```
//Naming the Columns of SPARK Dataframe into RDD
```

```
val columnnameap =
Seq("airportabbrev", "airportfullname", "sourcecity", "scityabbrev")
```

```
//Renaming the columns of SPARK Dataframe with the help of SPARK RDD passed
as an argument
```

```
val secondsapdf = apdf.toDF(columnnameap: _*)
```

```
//Naming the Columns of SPARK Dataframe into RDD
```

```
val columnnameap = Seq("apabbrev", "apfullname", "destcity", "dcityabbrev")
```

```
//Renaming the columns of SPARK Dataframe with the help of SPARK RDD
```

```
val seconddapdf = apdf.toDF(columnnameap: _*)
```

```
-----
//Conversion of airlines RDD into Dataset
```

```
val asds = airlines.toDS()
```

```

//Conversion of airlines Dataset to SPARK Dataframe
val aldf = asds.toDF()

//Naming the Columns of SPARK Dataframe into RDD
val columnnameal = Seq("airlineabbrev","airlinefullname")

//Renaming the columns of SPARK Dataframe with the help of SPARK RDD
val secondaldf = aldf.toDF(columnnameal:_* )
-----

//Broadcast Join(Inner Join in this scenario) of fact table flight and
broadcast table airports for destination city
val
firstjointdf=firstddf.join(broadcast(seconddapdf),firstddf("airportabbrev")
=== seconddapdf("apabbrev"),"inner")

//Display the Schema of Spark Dataframe generated broadcast join
firstjointdf.printSchema()

//Broadcast Join(Inner Join in this scenario) of fact table flight and
broadcast table airports for source city

val
secondjointdf=firstjointdf.join(broadcast(secondsapdf),firstjointdf("source
abbrev") === secondsapdf("airportabbrev"),"inner")

//Display the Schema of Spark Dataframe generated broadcast join
secondjointdf.printSchema()

//Broadcast Join(Inner Join in this scenario) of fact table flight and
broadcast table airlines
val thirdjointdf =
secondjointdf.join(broadcast(secondaldf),secondjointdf("airlineabbrev") ===
secondaldf("airlineabbrev"),"inner")

//Display the Schema of Spark Dataframe generated broadcast join
thirdjointdf.printSchema()

//Display the final output of broadcast join of fact table and dimension
tables
thirdjointdf.select(
$"sourcecity",
$"destcity",
$"airlinefullname",
$"track",
$"depttime"
).show()

```

Screenshots:

```

scala> import org.apache.spark.sql._
import org.apache.spark.sql._

scala> import sqlContext.implicits._
import sqlContext.implicits._

scala> val flights = sc.parallelize(List(
  ("SEA", "JFK", "DL", "418", "7:00"),
  ("SFO", "LAX", "AA", "1250", "7:05"),
  ("SFO", "JFK", "VX", "12", "7:05"),
  ("JFK", "LAX", "DL", "424", "7:10"),
  ("LAX", "SEA", "DL", "5737", "7:10")
))
flights: org.apache.spark.rdd.RDD[(String, String, String, String, String)] = ParallelCollectionRDD[93] at parallelize at <console>:45

scala> val airports = sc.parallelize(List(
  ("JFK", "John F. Kennedy International Airport", "New York", "NY"),
  ("LAX", "Los Angeles International Airport", "Los Angeles", "CA"),
  ("SEA", "Seattle-Tacoma International Airport", "Seattle", "WA"),
  ("SFO", "San Francisco International Airport", "San Francisco", "CA")
))
airports: org.apache.spark.rdd.RDD[(String, String, String, String)] = ParallelCollectionRDD[94] at parallelize at <console>:45

scala> val airlines = sc.parallelize(List(
  ("AA", "American Airlines"),
  ("DL", "Delta Airlines"),
  ("VX", "Virgin America")
))
airlines: org.apache.spark.rdd.RDD[(String, String)] = ParallelCollectionRDD[95] at parallelize at <console>:45

scala> val fds = flights.toDS()
fds: org.apache.spark.sql.Dataset[(String, String, String, String, String)] = [_1: string, _2: string, _3: string, _4: string, _5: string]

scala> val firstddf = fds.toDF()
firstddf: org.apache.spark.sql.DataFrame = [_1: string, _2: string, _3: string, _4: string, _5: string]

scala> val columnnamef = Seq("sourceabbrev", "airportabbrev", "airlineabbrev", "track", "depttime")
columnnamef: Seq[String] = List(sourceabbrev, airportabbrev, airlineabbrev, track, depttime)

scala> val firstddf = firstddf.toDF(columnnamef: *)
firstddf: org.apache.spark.sql.DataFrame = [sourceabbrev: string, airportabbrev: string, airlineabbrev: string, track: string, depttime: string]

scala> val apds = airports.toDS()
apds: org.apache.spark.sql.Dataset[(String, String, String, String)] = [_1: string, _2: string, _3: string, _4: string]

scala> val apdf = apds.toDF()
apdf: org.apache.spark.sql.DataFrame = [_1: string, _2: string, _3: string, _4: string]

scala> val columnnameap = Seq("airportabbrev", "airportfullname", "sourcecity", "scityabbrev")
columnnameap: Seq[String] = List(airportabbrev, airportfullname, sourcecity, scityabbrev)

scala> val secondsapdf = apdf.toDF(columnnameap: *)
secondsapdf: org.apache.spark.sql.DataFrame = [airportabbrev: string, airportfullname: string, sourcecity: string, scityabbrev: string]

scala> val columnnameap = Seq("apabbrev", "apfullname", "destcity", "dcityabbrev")
columnnameap: Seq[String] = List(apabbrev, apfullname, destcity, dcityabbrev)

scala> val seconddapdf = apdf.toDF(columnnameap: *)
seconddapdf: org.apache.spark.sql.DataFrame = [apabbrev: string, apfullname: string, destcity: string, dcityabbrev: string]

scala> val fds = airlines.toDS()
fds: org.apache.spark.sql.Dataset[(String, String)] = [_1: string, _2: string]

scala> val aldf = fds.toDF()
aldf: org.apache.spark.sql.DataFrame = [_1: string, _2: string]

scala> val columnnameal = Seq("airlineabbrev", "airlinefullname")
columnnameal: Seq[String] = List(airlineabbrev, airlinefullname)

scala> val secondaldf = aldf.toDF(columnnameal: *)
secondaldf: org.apache.spark.sql.DataFrame = [airlineabbrev: string, airlinefullname: string]

scala> val firstjointdf=firstddf.join(broadcast(seconddapdf),firstddf("airportabbrev") === seconddapdf("apabbrev"),"inner")
firstjointdf: org.apache.spark.sql.DataFrame = [sourceabbrev: string, airportabbrev: string, airlineabbrev: string, track: string, depttime: string, apabbrev: string, apfullname: string, destcity: string, dcityabbrev: string]

scala> firstjointdf.printSchema()
root
|-- sourceabbrev: string (nullable = true)
|-- airportabbrev: string (nullable = true)
|-- airlineabbrev: string (nullable = true)
|-- track: string (nullable = true)
|-- depttime: string (nullable = true)
|-- apabbrev: string (nullable = true)
|-- apfullname: string (nullable = true)
|-- destcity: string (nullable = true)
|-- dcityabbrev: string (nullable = true)

scala> val secondjointdf=firstjointdf.join(broadcast(secondsapdf),firstjointdf("sourceabbrev") === secondsapdf("airportabbrev"),"inner")
secondjointdf: org.apache.spark.sql.DataFrame = [sourceabbrev: string, airportabbrev: string, airlineabbrev: string, track: string, depttime: string, apabbrev: string, apfullname: string, destcity: string, dcityabbrev: string]

```

```

ring, dcityabbrev: string, airportabbrev: string, airportfullname: string, sourcecity: string, scity
abbrev: string]

scala> secondjointdf.printSchema()
root
|-- sourceabbrev: string (nullable = true)
|-- airportabbrev: string (nullable = true)
|-- airlineabbrev: string (nullable = true)
|-- track: string (nullable = true)
|-- depttime: string (nullable = true)
|-- apabbrev: string (nullable = true)
|-- apfullname: string (nullable = true)
|-- destcity: string (nullable = true)
|-- dcityabbrev: string (nullable = true)
|-- airportabbrev: string (nullable = true)
|-- airportfullname: string (nullable = true)
|-- sourcecity: string (nullable = true)
|-- scityabbrev: string (nullable = true)

scala> val thirdjointdf = secondjointdf.join(broadcast(secondaldf),secondjointdf("airlineabbrev") ==
=
| secondaldf("airlineabbrev"),"inner")
thirdjointdf: org.apache.spark.sql.DataFrame = [sourceabbrev: string, airportabbrev: string, airline
abbrev: string, track: string, depttime: string, apabbrev: string, apfullname: string, destcity: str
ing, dcityabbrev: string, airportabbrev: string, airportfullname: string, sourcecity: string, scitya
bbrev: string, airlineabbrev: string, airlinefullname: string]

```

```

scala> thirdjointdf.printSchema()
root
|-- sourceabbrev: string (nullable = true)
|-- airportabbrev: string (nullable = true)
|-- airlineabbrev: string (nullable = true)
|-- track: string (nullable = true)
|-- depttime: string (nullable = true)
|-- apabbrev: string (nullable = true)
|-- apfullname: string (nullable = true)
|-- destcity: string (nullable = true)
|-- dcityabbrev: string (nullable = true)
|-- airportabbrev: string (nullable = true)
|-- airportfullname: string (nullable = true)
|-- sourcecity: string (nullable = true)
|-- scityabbrev: string (nullable = true)
|-- airlineabbrev: string (nullable = true)
|-- airlinefullname: string (nullable = true)

```

Final Output:

```

scala> thirdjointdf.select(
|   $"sourcecity",
|   $"destcity",
|   $"airlinefullname",
|   $"track",
|   $"depttime"
| ).show()

```

sourcecity	destcity	airlinefullname	track	depttime
Los Angeles	Seattle	Delta Airlines	5737	7:10
New York	Los Angeles	Delta Airlines	424	7:10
Seattle	New York	Delta Airlines	418	7:00
San Francisco	Los Angeles	American Airlines	1250	7:05
San Francisco	New York	Virgin America	12	7:05